



# Data Mining with R part 1

Veerasak Kritsanapraphan

March 10, 2017

# Slide and Sample Data

[https://github.com/vkrit/chula\\_datamining](https://github.com/vkrit/chula_datamining).



# Agenda

1. Overview and data visualization
2. Data Preparation
3. Predictive Data Mining
  - Decision Tree
  - K-Nearest Neighbor
  - Naive Bayes Classifier
  - Neural Network

# Overview

Predictive Data Mining : two phases of processing

1. Training Phase : Create a model from training data
2. Predicting phase (Testing) : Deploy the model to production and use that to predict the future outcome

# Data

Iris Data Set from UCI Machine Learning Repository  
<https://archive.ics.uci.edu/ml/datasets/iris>

# Iris Data - Attribute Information

Column	Data Description
1	Sepal Length in cm
2	Sepal Width in cm
3	Petal Length in cm
4	Petal Width in cm
5	Classes

---

# Iris Data - Class (Label)

Column

Data

Class

- Iris Setosa
  - Iris Versicolor
  - Iris Verginica
-

# Getting Data

```
iris <- read.csv("iris.data.csv", header = TRUE)
summary(iris)
```

```
##   Sepal.Length   Sepal.Width   Petal.Length   Petal.Width
##   Min.    :4.300   Min.    :2.000   Min.    :1.000   Min.    :0.100
##   1st Qu.:5.100   1st Qu.:2.800   1st Qu.:1.600   1st Qu.:0.300
##   Median :5.800   Median :3.000   Median :4.350   Median :1.300
##   Mean    :5.843   Mean    :3.054   Mean    :3.759   Mean    :1.199
##   3rd Qu.:6.400   3rd Qu.:3.300   3rd Qu.:5.100   3rd Qu.:1.800
##   Max.    :7.900   Max.    :4.400   Max.    :6.900   Max.    :2.500
##           Species
##   Iris-setosa    :50
##   Iris-versicolor:50
##   Iris-virginica :50
##
##
##
```



# Exploring Data

```
nrow(iris)
```

```
## [1] 150
```

```
table(iris$Species)
```

```
##
```

```
##      Iris-setosa Iris-versicolor Iris-virginica
```

```
##              50              50              50
```

# Data Visualization

- Visualizing existing data is a very useful way to come up with ideas about what features should be included.
- "Dataframe" in R is a common way where data samples are organized in a tabular structure.

# Plot data according to their types


Let's start with numeric type first

*# Plot the histogram*

```
hist(iris$Sepal.Length, breaks = 10, prob = T)
```

*# Plot the density curve*

```
lines(density(iris$Sepal.Length))
```

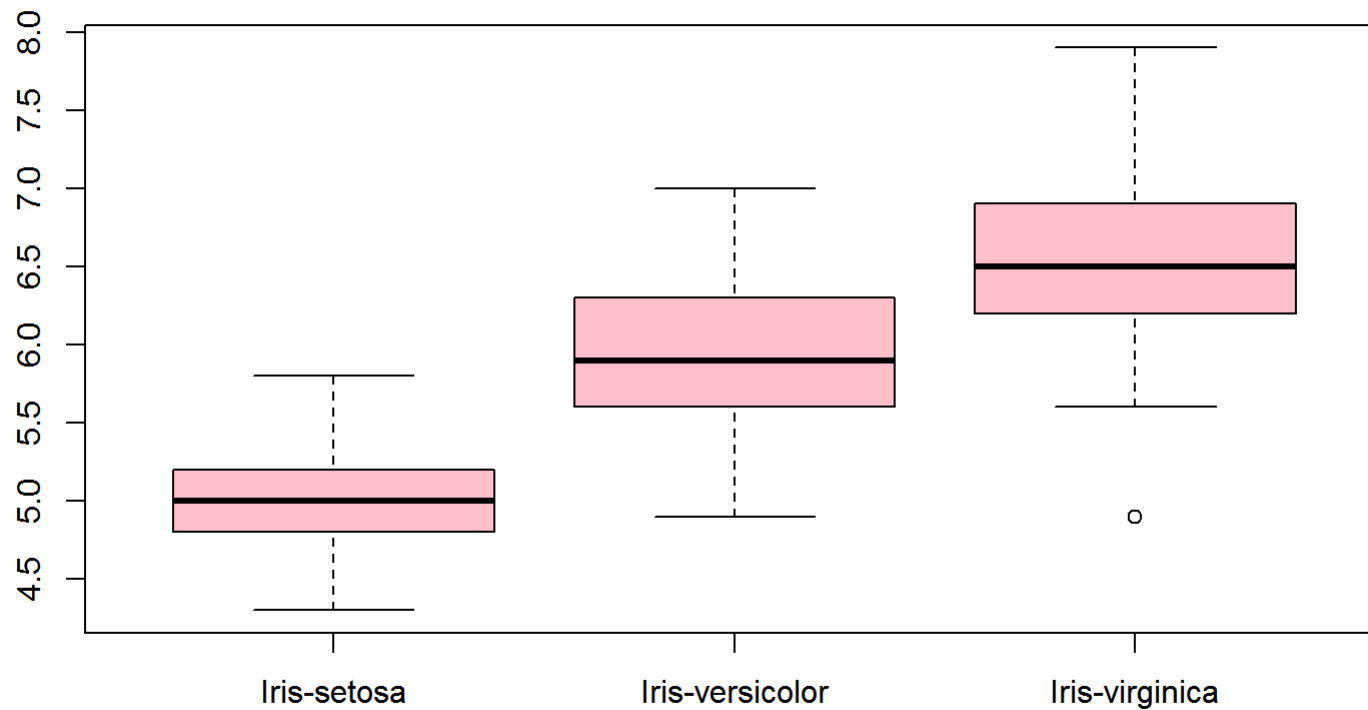


## Next data in categorical type

```
categories <- table(iris$Species)
barplot(categories, col = c("red", "white", "blue"))
```

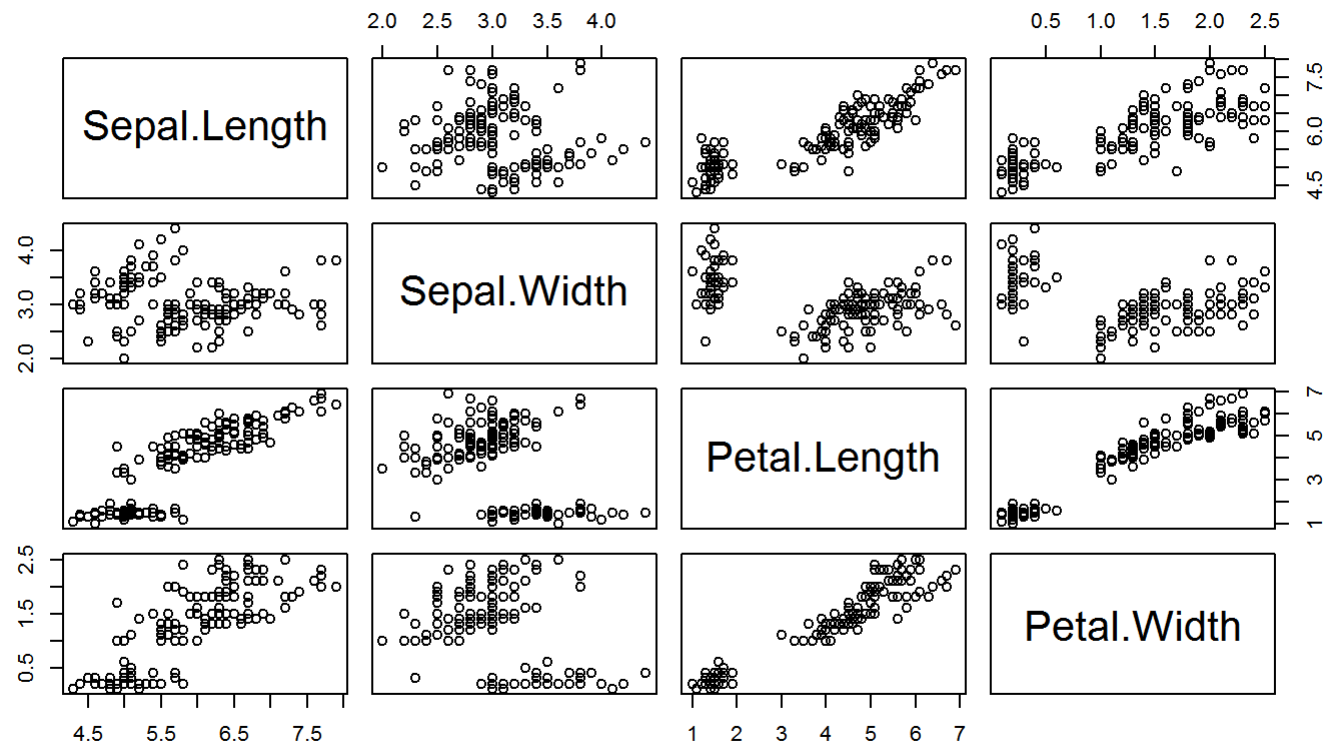
Let cross data between numeric and categorical

```
boxplot(Sepal.Length ~ Species, data = iris, col = c("pink"))
```



Next, let do correlation in charts using "pairs"

```
# Scatter plot for all pairs  
pairs(iris[, c(1, 2, 3, 4)])
```



## Compute the correlation matrix

```
# Compute the correlation matrix
```

```
cor(iris[, c(1, 2, 3, 4)])
```

```
##           Sepal.Length Sepal.Width Petal.Length Petal.Width
## Sepal.Length    1.0000000 -0.1093692    0.8717542    0.8179536
## Sepal.Width     -0.1093692    1.0000000   -0.4205161   -0.3565441
## Petal.Length     0.8717542   -0.4205161    1.0000000    0.9627571
## Petal.Width      0.8179536   -0.3565441    0.9627571    1.0000000
```

Draw regression line for suspect attribute to see relationship

```
plot(Petal.Width ~ Sepal.Length, data = iris)
model <- lm(Petal.Width ~ Sepal.Length, data = iris)
abline(model)
model2 <- lowess(iris$Petal.Width ~ iris$Sepal.Length)
lines(model2, col = "red")
```



# Preparing Training Data

At this step, the purpose is to transform the raw data in a form that can fit into the data mining model.

- Data sampling
- Data validation and handle missing data
- Normalize numeric value into a uniform range
- Compute aggregated value (a special case is to compute frequency counts)
- Expand categorical field to binary fields
- Discretize numeric value into categories
- Create derived fields from existing fields
- Reduce dimensionality
- Power and Log transformation

# Data Sampling

```
# Select 10 records out from iris with replacement  
index <- sample(1:nrow(iris), 10, replace = T)  
index
```

```
## [1] 90 107 115 62 26 16 1 116 87 60
```

```
# Subset iris to irissample from index
irissample <- iris[index, ]
irissample
```

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
## 90	5.5	2.5	4.0	1.3	Iris-versicolor
## 107	4.9	2.5	4.5	1.7	Iris-virginica
## 115	5.8	2.8	5.1	2.4	Iris-virginica
## 62	5.9	3.0	4.2	1.5	Iris-versicolor
## 26	5.0	3.0	1.6	0.2	Iris-setosa
## 16	5.7	4.4	1.5	0.4	Iris-setosa
## 1	5.1	3.5	1.4	0.2	Iris-setosa
## 116	6.4	3.2	5.3	2.3	Iris-virginica
## 87	6.7	3.1	4.7	1.5	Iris-versicolor
## 60	5.2	2.7	3.9	1.4	Iris-versicolor

# Impute missing data

- Discard the whole record
- Infer missing data with Average or Median

*# Create some missing data*

```
irissample[10, 1] <- NA
```

```
irissample
```

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
## 90	5.5	2.5	4.0	1.3	Iris-versicolor
## 107	4.9	2.5	4.5	1.7	Iris-virginica
## 115	5.8	2.8	5.1	2.4	Iris-virginica
## 62	5.9	3.0	4.2	1.5	Iris-versicolor
## 26	5.0	3.0	1.6	0.2	Iris-setosa
## 16	5.7	4.4	1.5	0.4	Iris-setosa
## 1	5.1	3.5	1.4	0.2	Iris-setosa
## 116	6.4	3.2	5.3	2.3	Iris-virginica
## 87	6.7	3.1	4.7	1.5	Iris-versicolor
## 60	NA	2.7	3.9	1.4	Iris-versicolor

*# Fix using Mean*

```
library(e1071)
```

```
fixiris1 <- impute(irissample[, 1:4], what = "mean")
```

```
fixiris1
```

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
## 90	5.500000	2.5	4.0	1.3
## 107	4.900000	2.5	4.5	1.7
## 115	5.800000	2.8	5.1	2.4
## 62	5.900000	3.0	4.2	1.5
## 26	5.000000	3.0	1.6	0.2
## 16	5.700000	4.4	1.5	0.4
## 1	5.100000	3.5	1.4	0.2
## 116	6.400000	3.2	5.3	2.3
## 87	6.700000	3.1	4.7	1.5
## 60	5.666667	2.7	3.9	1.4

```
# Fix using Median
```

```
library(e1071)
```

```
fixiris2 <- impute(irissample[, 1:4], what = "median")
```

```
fixiris2
```

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
## 90	5.5	2.5	4.0	1.3
## 107	4.9	2.5	4.5	1.7
## 115	5.8	2.8	5.1	2.4
## 62	5.9	3.0	4.2	1.5
## 26	5.0	3.0	1.6	0.2
## 16	5.7	4.4	1.5	0.4
## 1	5.1	3.5	1.4	0.2
## 116	6.4	3.2	5.3	2.3
## 87	6.7	3.1	4.7	1.5
## 60	5.7	2.7	3.9	1.4

# Normalize numeric value

*# Scale the columns  $x = \text{mean}(x) / \text{standard deviation}$*

```
scaleiris <- scale(iris[, 1:4])
```

```
head(scaleiris)
```

```
##      Sepal.Length Sepal.Width Petal.Length Petal.Width
## [1,]   -0.8976739    1.0286113   -1.336794   -1.308593
## [2,]   -1.1392005   -0.1245404   -1.336794   -1.308593
## [3,]   -1.3807271    0.3367203   -1.393470   -1.308593
## [4,]   -1.5014904    0.1060900   -1.280118   -1.308593
## [5,]   -1.0184372    1.2592416   -1.336794   -1.308593
## [6,]   -0.5353840    1.9511326   -1.166767   -1.046525
```

# Reduce dimension

There are two ways to reduce the number of input attributes.

1. Removing irrelevant input variables.
2. Removing redundant input variables.



```
# Use iris data set
```

```
cor(iris[, -5])
```

```
##           Sepal.Length Sepal.Width Petal.Length Petal.Width
## Sepal.Length    1.0000000 -0.1093692    0.8717542    0.8179536
## Sepal.Width     -0.1093692    1.0000000   -0.4205161   -0.3565441
## Petal.Length     0.8717542   -0.4205161    1.0000000    0.9627571
## Petal.Width      0.8179536   -0.3565441    0.9627571    1.0000000
```

```
# Compute PCA Notice that PC1 and PC2 covers most variation
```


```
pca <- prcomp(iris[, -5], scale = T)
```

```
summary(pca)
```

```
## Importance of components:
```

```
##           PC1      PC2      PC3      PC4
## Standard deviation    1.7061 0.9598 0.38387 0.14355
## Proportion of Variance 0.7277 0.2303 0.03684 0.00515
## Cumulative Proportion 0.7277 0.9580 0.99485 1.00000
```

```
# Plot PCA  
plot(pca)
```



pca\$rotation

##		PC1	PC2	PC3	PC4
##	Sepal.Length	0.5223716	-0.37231836	0.7210168	0.2619956
##	Sepal.Width	-0.2633549	-0.92555649	-0.2420329	-0.1241348
##	Petal.Length	0.5812540	-0.02109478	-0.1408923	-0.8011543
##	Petal.Width	0.5656110	-0.06541577	-0.6338014	0.5235463

```
predict(pca)[1:2, ]
```

```
##           PC1          PC2          PC3          PC4
## [1,] -2.256981 -0.5040154  0.1215362  0.02299628
## [2,] -2.079459  0.6532164  0.2264921  0.10286364
```

```
biplot(pca)
```

# Add derived attributes

```
iris2 <- transform(iris, ratio = round(Sepal.Length/Sepal.Width,  
  2))  
head(iris2)
```

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species	ratio
## 1	5.1	3.5	1.4	0.2	Iris-setosa	1.46
## 2	4.9	3.0	1.4	0.2	Iris-setosa	1.63
## 3	4.7	3.2	1.3	0.2	Iris-setosa	1.47
## 4	4.6	3.1	1.5	0.2	Iris-setosa	1.48
## 5	5.0	3.6	1.4	0.2	Iris-setosa	1.39
## 6	5.4	3.9	1.7	0.4	Iris-setosa	1.38

# Discretize numeric value into categories

```
# Equal width cuts
```

```
segments = 10
```

```
maxL <- max(iris$Petal.Length)
```

```
minL <- min(iris$Petal.Length)
```

```
theBreaks <- seq(minL, maxL, by = (maxL - minL)/segments)
```

```
cutPetalLength <- cut(iris$Petal.Length, breaks = theBreaks,  
  include.lowest = T)
```

```
newdata <- data.frame(origin.Petal.Len = iris$Petal.Length,  
  cut.Petal.Len = cutPetalLength)
```

```
head(newdata)
```

```
##   origin.Petal.Len cut.Petal.Len  
## 1             1.4      [1,1.59]  
## 2             1.4      [1,1.59]  
## 3             1.3      [1,1.59]  
## 4             1.5      [1,1.59]  
## 5             1.4      [1,1.59]  
## 6             1.7   (1.59,2.18]
```

*# Constant frequency / height*

```
myBreaks <- quantile(iris$Petal.Length, probs = seq(0, 1, 1/segments))
cutPetalLength2 <- cut(iris$Petal.Length, breaks = myBreaks,
  include.lowest = T)
newdata2 <- data.frame(orig.Petal.Len = iris$Petal.Length,
  cut.Petal.Len = cutPetalLength2)
head(newdata2)
```

```
##   orig.Petal.Len cut.Petal.Len
## 1           1.4      [1,1.4]
## 2           1.4      [1,1.4]
## 3           1.3      [1,1.4]
## 4           1.5    (1.4,1.5]
## 5           1.4      [1,1.4]
## 6           1.7    (1.7,3.9]
```

# Binarize categorical attributes

```
cat <- levels(iris$Species)
binarize <- function(x) {
  return(iris$Species == x)
}
newcols <- sapply(cat, binarize)
colnames(newcols) <- cat
data <- cbind(iris[, c("Species")], newcols)
```



# Binarize categorical attributes (cont)

```
data[45:55, ]
```

```
##           Iris-setosa Iris-versicolor Iris-virginica
## [1,] 1             1                 0                0
## [2,] 1             1                 0                0
## [3,] 1             1                 0                0
## [4,] 1             1                 0                0
## [5,] 1             1                 0                0
## [6,] 1             1                 0                0
## [7,] 2             0                 1                0
## [8,] 2             0                 1                0
## [9,] 2             0                 1                0
## [10,] 2            0                 1                0
## [11,] 2            0                 1                0
```

# Data Mining Techniques

# Iris Data Preparation

```
# sample iris into 2 sets (training, testing) with 70%/30%  
set.seed(1234)  
ind <- sample(2, nrow(iris), replace = TRUE, prob = c(0.7,  
  0.3))  
trainData <- iris[ind == 1, ]  
testData <- iris[ind == 2, ]
```

# Decision Tree

Libray name -> party

Function name -> ctree

Installation

```
install.packages("party")
```

# Decision Tree - Create Model

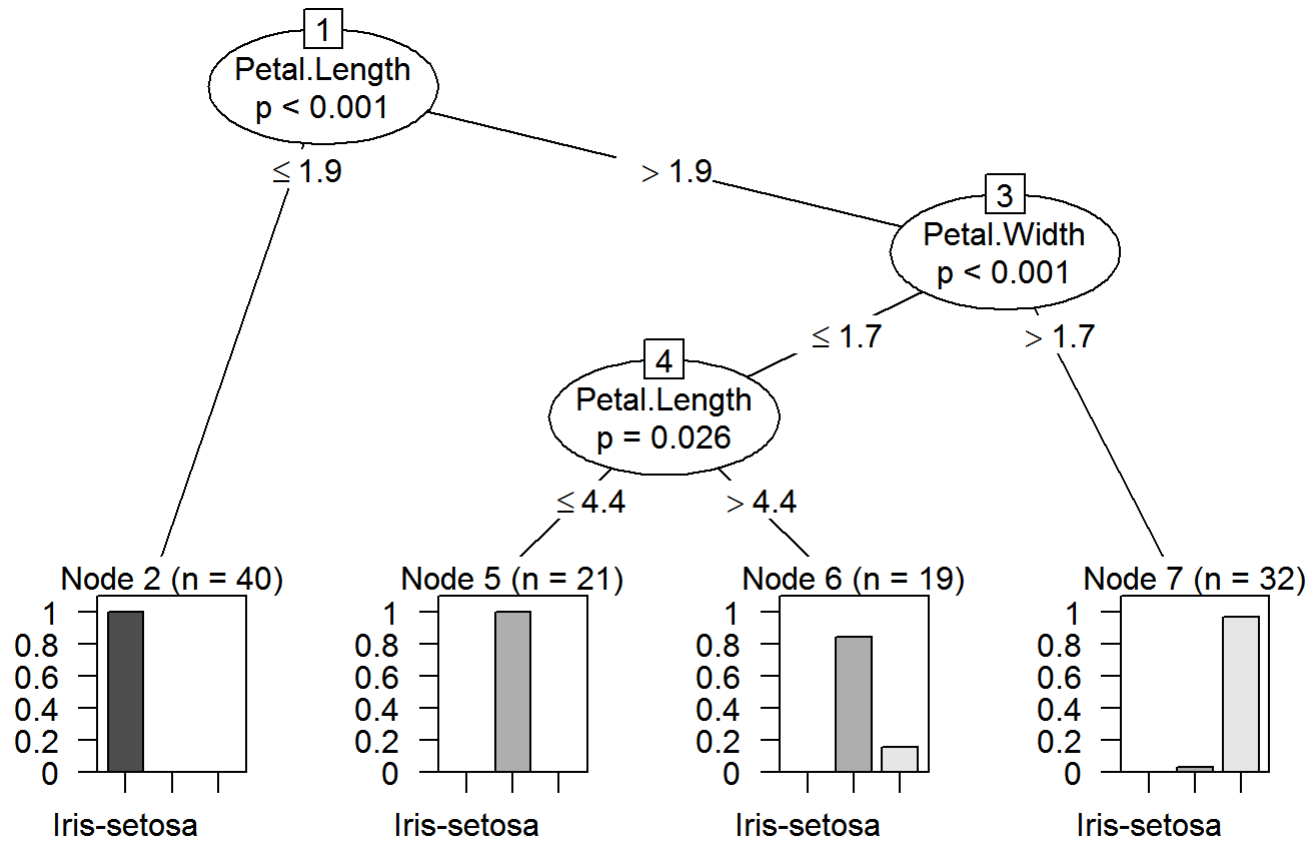
```
library(party)
myFormula <- Species ~ Sepal.Length + Sepal.Width + Petal.Length +
  Petal.Width
iris_ctree <- ctree(myFormula, data = trainData)
# Check the prediction
table(predict(iris_ctree), trainData$Species)
```

```
##
##           Iris-setosa Iris-versicolor Iris-virginica
## Iris-setosa           40              0              0
## Iris-versicolor        0              37              3
## Iris-virginica         0              1             31
```

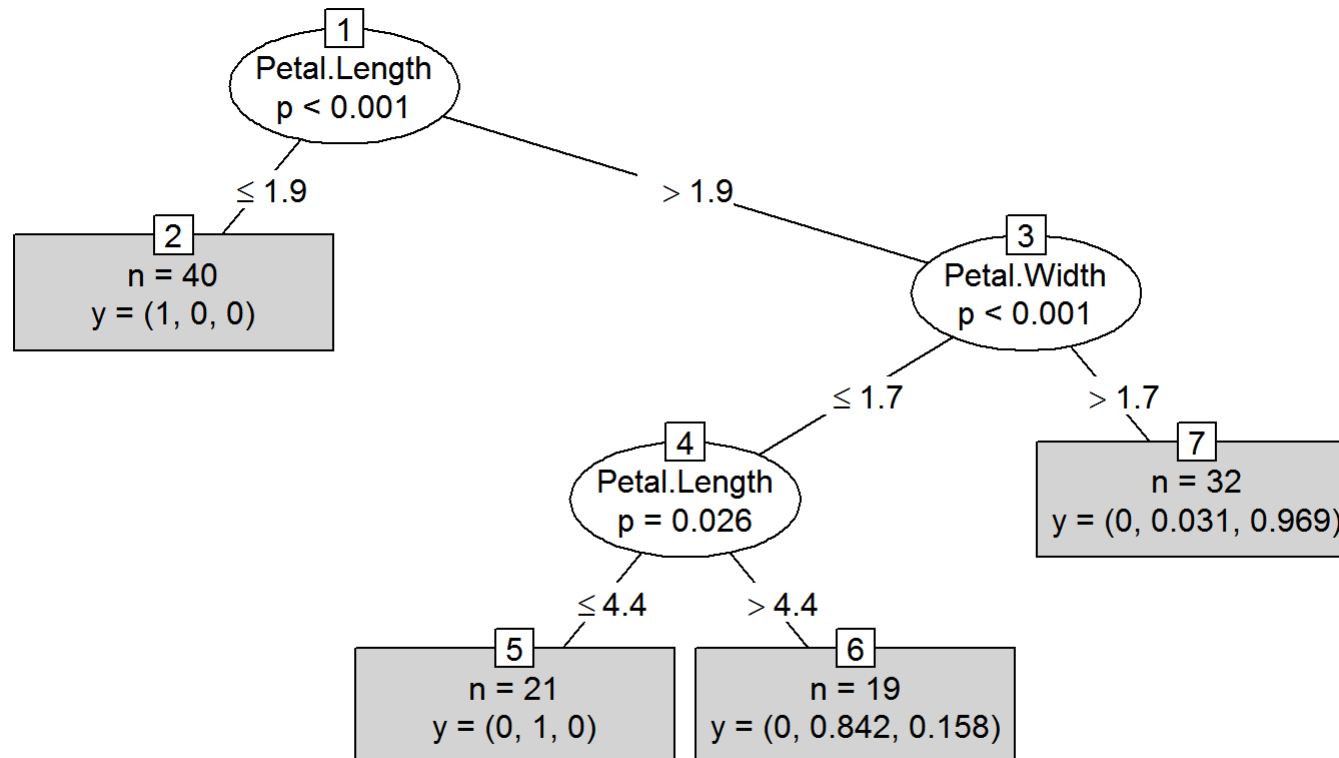
```
print(iris_ctree)

##
##  Conditional inference tree with 4 terminal nodes
##
## Response:  Species
## Inputs:  Sepal.Length, Sepal.Width, Petal.Length, Petal.Width
## Number of observations:  112
##
## 1) Petal.Length <= 1.9; criterion = 1, statistic = 104.637
##   2)* weights = 40
## 1) Petal.Length > 1.9
##   3) Petal.Width <= 1.7; criterion = 1, statistic = 48.939
##     4) Petal.Length <= 4.4; criterion = 0.974, statistic = 7.397
##       5)* weights = 21
##       4) Petal.Length > 4.4
##         6)* weights = 19
##     3) Petal.Width > 1.7
##       7)* weights = 32
```

```
plot(iris_ctree)
```



```
plot(iris_ctree, type = "simple")
```





# Decision Tree - Prediction

```
# Predict on test data
```

```
testPred <- predict(iris_ctree, newdata = testData)  
table(testPred, testData$Species)
```

```
##
```

```
## testPred      Iris-setosa Iris-versicolor Iris-virginica  
##  Iris-setosa           10              0              0  
##  Iris-versicolor        0             12              2  
##  Iris-virginica         0              0             14
```

# K-Nearest Neighbor Classification (k-NN)

Library -> class

Function -> knn

Installation

```
install.packages("class")
```

# K-NN

```
library(class)
train_input <- as.matrix(trainData[, -5])
train_output <- as.vector(trainData[, 5])
test_input <- as.matrix(testData[, -5])
prediction <- knn(train_input, test_input, train_output, k = 5)
table(prediction, testData$Species)
```

```
##
## prediction      Iris-setosa Iris-versicolor Iris-virginica
##  Iris-setosa           10             0             0
##  Iris-versicolor        0             12             0
##  Iris-virginica         0             0             16
```

# Naive Bayes Classifier

Library -> e1071

Function -> naiveBayes

Installation

```
install.packages("e1071")
```

# Naive Bayes

```
library(e1071)
# can handle both categorical and numeric input but output
# must be categorical
model <- naiveBayes(Species ~ ., data = trainData)
prediction <- predict(model, testData[, -5])
table(prediction, testData[, 5])
```

```
##
## prediction      Iris-setosa Iris-versicolor Iris-virginica
## Iris-setosa      10           0              0
## Iris-versicolor   0          12              2
## Iris-virginica    0           0             14
```

# Neural Network

Library -> neuralnet

Function -> neuralnet

Installation

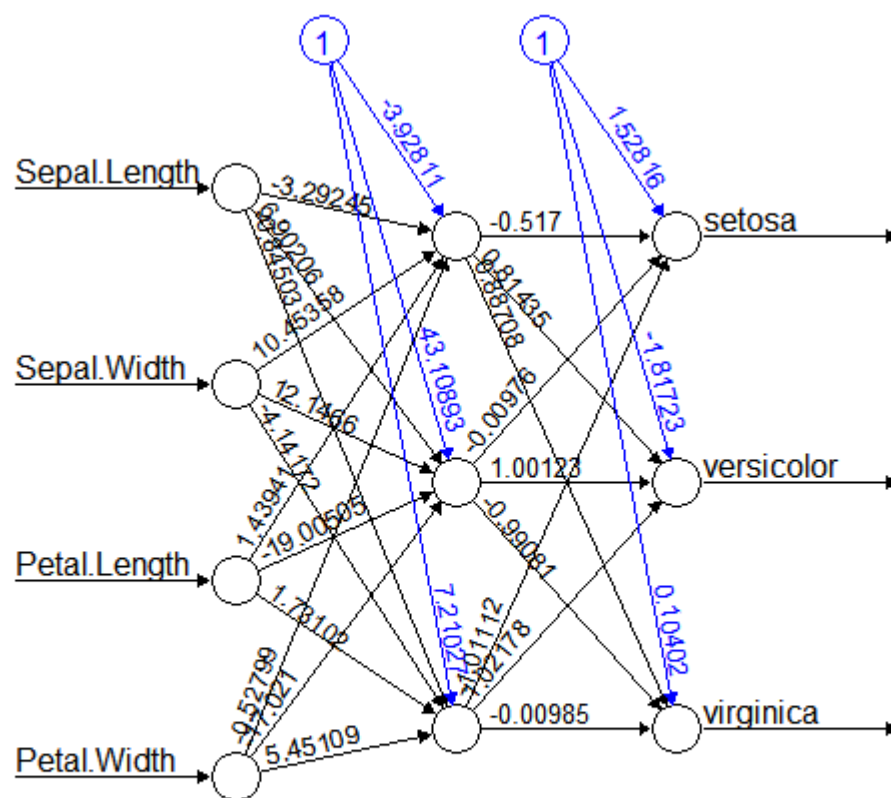
```
install.packages('neuralnet')
```

# Neural Network

```
library(neuralnet)
nnet_trainData <- trainData
# Binarize the categorical output
nnet_trainData <- cbind(nnet_trainData, trainData$Species ==
  "Iris-setosa")
nnet_trainData <- cbind(nnet_trainData, trainData$Species ==
  "Iris-versicolor")
nnet_trainData <- cbind(nnet_trainData, trainData$Species ==
  "Iris-virginica")
names(nnet_trainData)[6] <- "setosa"
names(nnet_trainData)[7] <- "versicolor"
names(nnet_trainData)[8] <- "virginica"
nn <- neuralnet(setosa + versicolor + virginica ~ Sepal.Length +
  Sepal.Width + Petal.Length + Petal.Width, data = nnet_trainData,
  hidden = c(3))
```

# Neural Network - Plot Model

```
plot(nn)
```



Error: 1.237788 Steps: 24462



nn

# Neural Network - Prediction

```

mypredict <- compute(nn, testData[-5])$net.result
# put multiple binary output to categorical output
maxidx <- function(arr) {
  return(which(arr == max(arr)))
}
idx <- apply(mypredict, c(1), maxidx)
prediction <- c("Iris-setosa", "Iris-versicolor", "Iris-virginica")[idx]
table(prediction, testData$Species)

```

```

##
## prediction      Iris-setosa Iris-versicolor Iris-virginica
##  Iris-setosa           10             0             0
##  Iris-versicolor        0             12             0
##  Iris-virginica         0             0             16

```

# End of Part 1

