



Data Mining Part 2

Veerasak Kritsanapraphan

4/6/2017

Agenda

- Data Mining Techniques using R
- Predictive Modeling Performance
- Clustering
 - K-Means Clustering
 - Hierarchical Clustering
- Association Rules (Market Basket Analysis)
- Multi-model Learning
- Text-Mining

Slide and Sample Data

https://github.com/vkrit/chula_datamining.



Prepare Data

```
# Prepare iris
set.seed(567)
ind <- sample(2, nrow(iris), replace = TRUE, prob = c(0.7,
  0.3))
traindata <- iris[ind == 1, ]
testdata <- iris[ind == 2, ]
table(traindata$Species)
```

```
##
##      setosa versicolor virginica
##      35          37          35
```

Predictive Modeling Performance

		True class		Measures
		Positive	Negative	
Predicted class	Positive	True positive <i>TP</i>	False positive <i>FP</i>	Positive predictive value (PPV) $\frac{TP}{TP+FP}$
	Negative	False negative <i>FN</i>	True negative <i>TN</i>	Negative predictive value (NPV) $\frac{TN}{FN+TN}$
Measures		Sensitivity $\frac{TP}{TP+FN}$	Specificity $\frac{TN}{FP+TN}$	Accuracy $\frac{TP+TN}{TP+FP+FN+TN}$

5/51

Create Decision Tree Model

```
library(party)
```

```
myFormula <- Species ~ Sepal.Length + Sepal.Width + Petal.Length +  
  Petal.Width
```

```
iris_ctree <- ctree(myFormula, data = traindata)
```

Create Confusion Matrix from Training Data

```
library(caret)
trainPred = predict(iris_ctree, traindata)
confusionMatrix(traindata$Species, trainPred)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
## Prediction  setosa versicolor virginica
##   setosa      35          0          0
## versicolor    0          36          1
## virginica     0           3         32
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
##           Accuracy : 0.9626
```

7/51

Confusion Matrix

Confusion Matrix and Statistics				
Prediction	Reference			
	Iris-setosa	Iris-versicolor	Iris-virginica	
Iris-setosa	35	0	0	
Iris-versicolor	0	36	1	
Iris-virginica	0	3	32	
Overall Statistics				
Accuracy : 0.9626				
95% CI : (0.907, 0.9897)				
No Information Rate : 0.3645				
P-Value [Acc > NIR] : < 2.2e-16				
Kappa : 0.9439				
McNemar's Test P-Value : NA				
Statistics by Class:				
	Class: Iris-setosa	Class: Iris-versicolor	Class: Iris-virginica	
Sensitivity	1.0000	0.9231	0.9697	
Specificity	1.0000	0.9853	0.9595	
Pos Pred Value	1.0000	0.9730	0.9143	
Neg Pred Value	1.0000	0.9571	0.9861	
Prevalence	0.3271	0.3645	0.3084	
Detection Rate	0.3271	0.3364	0.2991	
Detection Prevalence	0.3271	0.3458	0.3271	
Balanced Accuracy	1.0000	0.9542	0.9646	

K-Means Clustering

1. Pick an initial set of K centroids (this can be random or any other means)
2. For each data point, assign it to the member of the closest centroid according to the given distance function
3. Adjust the centroid position as the mean of all its assigned member data points. Go back to (2) until the membership isn't change and centroid position is stable.
4. Output the centroids

K-Means (cont.)

```
library(stats)
set.seed(101)
km <- kmeans(iris[, 1:4], 3)
plot(iris[, 1], iris[, 2], col = km$cluster)
points(km$centers[, c(1, 2)], col = 1:2, pch = 19, cex = 2)
```

K-Means (cont.)

```
table(km$cluster, iris$Species)
```

```
##  
##      setosa versicolor virginica  
##  1         0           48         14  
##  2        50           0          0  
##  3         0           2         36
```

K-Means (second round)

```
set.seed(900)
km <- kmeans(iris[, 1:4], 3)
plot(iris[, 1], iris[, 2], col = km$cluster)
points(km$centers[, c(1, 2)], col = 1:3, pch = 19, cex = 2)
```

K-Means (second round - cont.)

```
##  
##      setosa versicolor virginica  
##  1         0          46         50  
##  2        17           4          0  
##  3        33           0          0
```


Hierarchical Clustering

1. Compute distance between every pairs of point/cluster
 - Distance between point is just using the distance function.
 - Compute distance between point A to cluster B may involve many choices (such as the min/max/avg distance between the point A and points in the cluster B)
 - Compute distance between cluster A to cluster B may first compute distance of all points pairs (one from cluster A and the the other from cluster B) and then pick either min/max/avg of these pairs.
2. Combine the two closest point/cluster into a cluster. Go back to (1) until only one big cluster remains

Hierarchical Clustering

```
set.seed(101)
sampleiris <- iris[sample(1:150, 40), ] # get samples from iris dataset
# each observation has 4 variables, ie, they are
# interpreted as 4-D points
distance <- dist(sampleiris[, -5], method = "euclidean")
cluster <- hclust(distance, method = "average")
```

```
plot(cluster, hang = -1, label = sampleiris$Species)
```



Prune the result tree to 3 groups

```
group.3 <- cutree(cluster, k = 3) # prune the tree  
table(group.3, sampleiris$Species)
```

```
##  
## group.3 setosa versicolor virginica  
##      1      0      15      9  
##      2     13      0      0  
##      3      0      0      3
```

Plot cluster by column 1 and 2

```
par(mfrow = c(1, 2))  
plot(sampleiris[, c(1, 2)], col = group.3, pch = 19, cex = 1,  
      main = "3 clusters")  
plot(sampleiris[, c(1, 2)], col = sampleiris$Species, pch = 19,  
      cex = 1, main = "real clusters")
```

Association Rules (Market Basket Analysis)



Support: The rule holds with support sup in T (the transaction data set) if $\text{sup} \%$ of transactions contain $X \rightarrow Y$.

Confidence: The rule holds in T with confidence conf if $\text{conf}\%$ of transactions that contain X also contain Y .

Lift : The Lift of the rule is $X \Rightarrow Y$ is the confidence of the rule divided by the expected confidence, assuming that the item sets are independent.

Apriori Algorithm

Load the libraries

```
library(registry)
```

```
library(Matrix)
```

```
library(arules)
```

```
library(arulesViz)
```

```
library(datasets)
```

Load the data set

```
data(Groceries)
```

Data Format

Explore Groceries Data

```
# Create an item frequency plot for the top 20 items  
itemFrequencyPlot(Groceries, topN = 20, type = "absolute")
```

Create Association Rules

```
rules <- apriori(Groceries, parameter = list(supp = 0.001,  
      conf = 0.8))
```

```
## Apriori
```

```
##
```

```
## Parameter specification:
```

```
## confidence minval  smax  arem  aval originalSupport  support minlen maxlen
```

```
##           0.8     0.1    1 none  FALSE              TRUE   0.001     1     10
```

```
## target    ext
```

```
## rules FALSE
```

```
##
```

```
## Algorithmic control:
```

```
## filter tree heap memopt load sort verbose
```

```
##    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
```

```
##
```

23/51

Plot Rules

```
plot(rules)
```

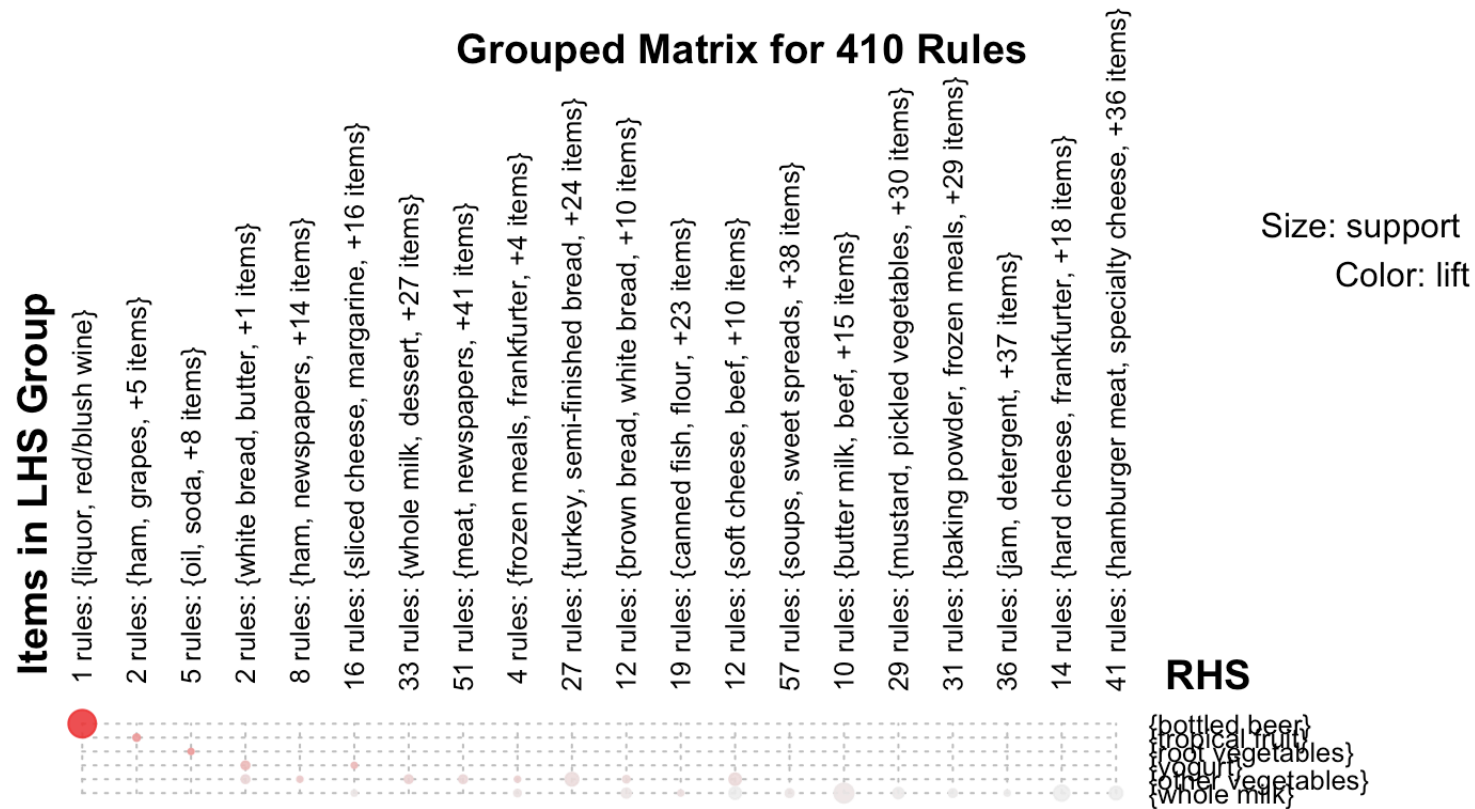

Sort Rules

```
rules <- sort(rules, by = "confidence", decreasing = TRUE)
inspect(rules[1:5])
```

##	lhs	rhs	support	confidence	lift
## 1	{rice, sugar}	=> {whole milk}	0.0012	1	3.9
## 2	{canned fish, hygiene articles}	=> {whole milk}	0.0011	1	3.9
## 3	{root vegetables, butter, rice}	=> {whole milk}	0.0010	1	3.9
## 4	{root vegetables, whipped/sour cream, flour}	=> {whole milk}	0.0017	1	3.9
## 5	{butter,				

25/51

```
plot(rules, method = "grouped")
```



26/51

Change to have limit association in one rule

```
rules <- apriori(Groceries, parameter = list(supp = 0.001,  
      conf = 0.8, maxlen = 3))
```

```
## Apriori
```

```
##
```

```
## Parameter specification:
```

```
## confidence minval  smax  arem  aval originalSupport  support minlen maxlen
```

```
##           0.8     0.1    1 none  FALSE                TRUE   0.001     1     3
```

```
## target    ext
```

```
## rules FALSE
```

```
##
```

```
## Algorithmic control:
```

```
## filter tree heap memopt load sort verbose
```

```
##    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
```


```
##
```

27/51

```
plot(rules, method = "graph")
```

Graph for 29 rules

size: support (0.001 - 0.003)
color: lift (3.131 - 11.235)



28/51

Rules pruned

```
subset.matrix <- is.subset(rules, rules)
subset.matrix[lower.tri(subset.matrix, diag = T)] <- NA
redundant <- colSums(subset.matrix, na.rm = T) >= 1
rules.pruned <- rules[!redundant]
rules <- rules.pruned
```

```
summary(rules)
```

```
## set of 29 rules
```

```
##
```

```
## rule length distribution (lhs + rhs):sizes
```

```
## 3
```

```
## 29
```

```
##
```

```
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
```

```
##           3         3         3         3         3         3
```

```
##
```

```
## summary of quality measures:
```

```
##      support      confidence      lift
```

```
## Min.      :0.00102  Min.      :0.80  Min.      : 3.1
```

```
## 1st Qu.:0.00112  1st Qu.:0.81  1st Qu.: 3.3
```

```
## Median :0.00122  Median :0.85  Median : 3.6
```

```
## Mean    :0.00147  Mean    :0.86  Mean    : 4.0
```

30/51

Targeting Items

- What are customers likely to buy before buying whole milk?
- What are customers likely to buy if they purchase whole milk?
- This essentially means we want to set either the Left Hand Side and Right Hand Side. This is not difficult to do with R!

Find whole milk's antecedents

```
rules <- apriori(data = Groceries, parameter = list(supp = 0.001,
  conf = 0.08), appearance = list(default = "lhs", rhs = "whole milk"),
  control = list(verbose = F))
rules <- sort(rules, decreasing = TRUE, by = "confidence")
inspect(rules[1:5])
```

##	lhs	rhs	support	confidence	lift
## 1	{rice,				
##	sugar}	=> {whole milk}	0.0012	1	3.9
## 2	{canned fish,				
##	hygiene articles}	=> {whole milk}	0.0011	1	3.9
## 3	{root vegetables,				
##	butter,				
##	rice}	=> {whole milk}	0.0010	1	3.9
## 4	{root vegetables,				

32/51

Likely to buy after buy whole milk

```
rules <- apriori(data = Groceries, parameter = list(supp = 0.001,
  conf = 0.15, minlen = 2), appearance = list(default = "rhs",
  lhs = "whole milk"), control = list(verbose = F))
rules <- sort(rules, decreasing = TRUE, by = "confidence")
inspect(rules[1:5])
```

##	lhs	rhs	support	confidence	lift
## 6	{whole milk}	=> {other vegetables}	0.075	0.29	1.5
## 5	{whole milk}	=> {rolls/buns}	0.057	0.22	1.2
## 4	{whole milk}	=> {yogurt}	0.056	0.22	1.6
## 2	{whole milk}	=> {root vegetables}	0.049	0.19	1.8
## 1	{whole milk}	=> {tropical fruit}	0.042	0.17	1.6

33/51

Bagging and Boosting using R

Multi-model Learning

Ensemble : Bagging

Random Forest

- Here is how such a system is trained; for some number of trees T :
- Sample N cases at random with replacement to create a subset of the data. The subset should be about 66% of the total set.
- At each node:
 - For some number m (see below), m predictor variables are selected at random from all the predictor variables
 - The predictor variable that provides the best split, according to some objective function, is used to do a binary split on that node.
 - At the next node, choose another m variables at random from all predictor variables and do the same.

36/51

Bagging

```
library(ggplot2)
library(randomForest)
# Train 500 trees, random selected attributes
model <- randomForest(Species ~ ., data = traindata, nTree = 500)
prediction <- predict(model, newdata = testdata, type = "class")
table(prediction, testdata$Species)
```

```
##
## prediction    setosa versicolor virginica
##   setosa         15          0          0
##   versicolor     0          13          2
##   virginica      0          0         13
```

37/51

Boosting

```
library(adabag)
iris.adaboost <- boosting(Species ~ ., data = traindata, boost = TRUE,
  mfinal = 5)
iris.adaboost
```

```
## $formula
## Species ~ .
##
## $trees
## $trees[[1]]
## n= 107
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
```

38/51

Plot variables important

```
barplot(iris.adaboost$imp[order(iris.adaboost$imp, decreasing = TRUE)],  
        ylim = c(0, 100), main = "Variables Relative Importance",  
        col = "lightblue")
```

Boosting (compare result)

```
table(iris.adaboost$class, traindata$Species, dnn = c("Predicted Class",  
  "Observed Class"))
```

```
##              Observed Class  
## Predicted Class setosa versicolor virginica  
##      setosa      35           0           0  
##      versicolor  0           35           0  
##      virginica   0           2          35
```


Text Mining

- Get Text Mining Library

```
# Needed <- c('tm', 'SnowballCC', 'RColorBrewer',  
# 'wordcloud', 'biclust', 'igraph', 'fpc')  
# install.packages(Needed, dependencies = TRUE)
```

Load file (Shakespear's Plays)

```
TEXTFILE = "t8.shakespeare.txt"
if (!file.exists(TEXTFILE)) {
  download.file("https://ocw.mit.edu/ans7870/6/6.006/s08/lecturenotes/files/t8.shake
    destfile = TEXTFILE)
}
shakespeare = readLines(TEXTFILE)
length(shakespeare)

## [1] 124456

shakespeare = shakespeare[-(1:173)]
shakespeare = shakespeare[-(124195:length(shakespeare))]
shakespeare = paste(shakespeare, collapse = " ")
shakespeare = strsplit(shakespeare, "<[^>]*>")[[1]]
```

42/51

Text Mining

```
library(tm)
docs.vec <- VectorSource(shakespeare)
docs.corpus <- Corpus(docs.vec)
summary(docs.corpus)
```

##	Length	Class	Mode
## 1	2	PlainTextDocument	list
## 2	2	PlainTextDocument	list
## 3	2	PlainTextDocument	list
## 4	2	PlainTextDocument	list
## 5	2	PlainTextDocument	list
## 6	2	PlainTextDocument	list
## 7	2	PlainTextDocument	list
## 8	2	PlainTextDocument	list
## 9	2	PlainTextDocument	list

43/51

Text Mining Basic

Remove Punctuation

```
docs.corpus <- tm_map(docs.corpus, removePunctuation)
head(docs.corpus)
```

```
## <<SimpleCorpus>>
```

```
## Metadata:  corpus specific: 1, document level (indexed): 0
```

```
## Content:  documents: 6
```

Remove Number

```
docs.corpus <- tm_map(docs.corpus, removeNumbers)
```

```
docs.corpus <- tm_map(docs.corpus, tolower)
```

Remove Stopwords

```
docs.corpus <- tm_map(docs.corpus, removeWords, stopwords("english"))
```

44/51

Clean Data

```
# remove ing s, es  
library(SnowballC)  
docs.corpus <- tm_map(docs.corpus, stemDocument)  
docs.corpus <- tm_map(docs.corpus, stripWhitespace)
```

Step of Text Mining

Create Document Term Matrix

```
# Create Document Term Matrix
dtm <- DocumentTermMatrix(docs.corpus)
inspect(dtm[1:10, 1:10])

## <<DocumentTermMatrix (documents: 10, terms: 10)>>
## Non-/sparse entries: 19/81
## Sparsity           : 81%
## Maximal term length: 6
## Weighting           : term frequency (tf)
## Sample              :
##      Terms
## Docs accept addit agent allow alter altern appli aris asi associ
##      1         1         1         1         1         1         2         1         1         1         1
```

46/51

Create Term Document Matrix

```
# Create Term Document Matrix
```

```
tdm <- TermDocumentMatrix(docs.corpus)
```

```
inspect(tdm[1:10, 1:10])
```

```
## <<TermDocumentMatrix (terms: 10, documents: 10)>>
```

```
## Non-/sparse entries: 19/81
```

```
## Sparsity           : 81%
```

```
## Maximal term length: 6
```

```
## Weighting          : term frequency (tf)
```

```
## Sample             :
```

```
##           Docs
```

```
## Terms      1 10 2 3 4 5 6 7 8 9
```

```
##   accept  1  0 2 0 0 0 0 0 0 0
```

```
##   addit   1  0 2 0 0 1 0 1 0 0
```

```
##   agent   1  0 0 0 0 0 0 0 0 0
```

47/51

Explore Data

```
# Explore Data
```

```
freq <- colSums(as.matrix(dtm))  
length(freq)
```

```
## [1] 18786
```

```
ord <- order(freq)  
head(ord)
```

```
## [1] 9 11 13 14 15 20
```


Removing sparse terms

```
# Start by removing sparse terms:
```

```
TDM.common = removeSparseTerms(tdm, 0.1)  
dim(tdm)
```

```
## [1] 18786    219
```

```
dim(TDM.common)
```

```
## [1]    0 219
```

```
m = as.matrix(tdm)  
v = sort(rowSums(m), decreasing = TRUE)  
d <- data.frame(word = names(v), freq = v)  
head(d, 10)
```

49/51

Create Word Cloud

```
library(wordcloud)
set.seed(1234)
wordcloud(words = d$word, freq = d$freq, min.freq = 1, max.words = 200,
  random.order = FALSE, rot.per = 0.35, colors = brewer.pal(8,
    "Dark2"))
```

email : veerasak.kr568@cbs.chula.ac.th

Thank you